

Assignment 2

Name: Aaron Bengochea - ab6503

School: New York University – Tandon School of Engineering

Program: Masters of Science in Computer Science

Part 1 & 2: Deploying the Application on Minikube

1) We create the Dockerfile for the application

```
Dockerfile > ...
1  # Use an official Python runtime as a parent image
2  FROM python:3.9-slim
3
4  # Set the working directory in the container
5  WORKDIR /app
6
7  # Copy the requirements file into the container
8  COPY . /app
9
10 # Install any needed packages specified in requirements.txt
11 RUN pip install --no-cache-dir --upgrade pip && \
12     pip install --no-cache-dir -r requirements.txt
13
14
15 # Expose the port that your Flask app runs on
16 EXPOSE 3000
17
18 # Define environment variable for Flask
19 ENV FLASK_APP=app.py
20
21 # Run the Flask application
22 CMD ["flask", "run", "--host=0.0.0.0", "--port=3000"]
23
```

2) We create the Docker-compose.yaml file using both the flask-app and mongodb images to create a new combined image.

```

🐳 docker-compose.yml
1  version: '3.8'
2
   ▶Run All Services
3  services:
   ▶Run Service
4  web:
5      build: .
6      image: aaronbengo/cca2:latest
7      ports:
8          - "3000:3000"
9      depends_on:
10         - mongo
11      environment:
12         - FLASK_ENV=development
13         - MONGO_URI=mongodb://mongo:27017/todo_db
14         - PORT=3000
15
   ▶Run Service
16  mongo:
17      image: mongo:4.4
18      ports:
19          - "27017:27017"
20      volumes:
21          - mongo-data:/data/db
22
23  volumes:
24      mongo-data:
25

```

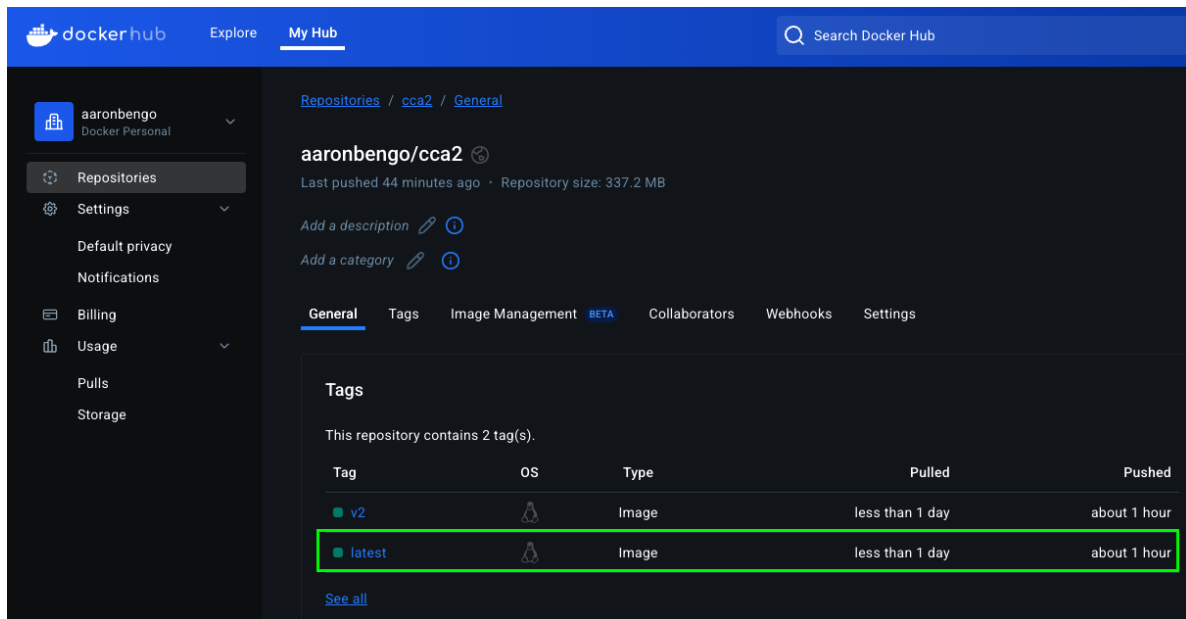
3) We push the finalized image to dockerhub using the following command in the CLI

```

docker buildx create --use docker buildx build --platform linux/amd64,linux/arm64 \
    -t aaronbengo/cca2:latest --push .

```

4) We observe that the image was successfully pushed to DockerHub with the name "aaronbengo/cca2:latest"



The screenshot shows the Docker Hub interface for the repository 'aaronbengo/cca2'. The left sidebar contains navigation links: 'aaronbengo Docker Personal', 'Repositories', 'Settings', 'Default privacy', 'Notifications', 'Billing', 'Usage', 'Pulls', and 'Storage'. The main content area shows the repository details, including the name 'aaronbengo/cca2', the last push time '44 minutes ago', and the repository size '337.2 MB'. Below this, there are links to 'Add a description' and 'Add a category'. The 'General' tab is selected, showing a table of tags. The table has columns for 'Tag', 'OS', 'Type', 'Pulled', and 'Pushed'. Two tags are listed: 'v2' and 'latest'. The 'latest' tag is highlighted with a green border. Below the table, there is a link to 'See all'.

Tag	OS	Type	Pulled	Pushed
v2	linux	Image	less than 1 day	about 1 hour
latest	linux	Image	less than 1 day	about 1 hour

Part 3: Deploying the Application on Minikube

1) Apply the deployment and services to minikube

```
→ cca2 git:(main) x kubectl apply -f flask-deployment.yaml
kubectl apply -f flask-service.yaml
kubectl apply -f mongodb-deployment.yaml
kubectl apply -f mongodb-service.yaml

deployment.apps/flask-app-deployment created
service/flask-app-service created
error: the path "mongodb-deployment.yaml" does not exist
error: the path "mongodb-service.yaml" does not exist
→ cca2 git:(main) x kubectl apply -f mongo-deployment.yaml
kubectl apply -f mongo-service.yaml

deployment.apps/mongo-deployment created
service/mongo created
```

2) We check deployments, services, and pods in order to verify that our deployment is working as expected. We find that our services are running as expected. We are currently running two replicas of our flask application. We can observe that there are two replicas in the "kubectl get deployments" command, and we can observe their pod names and their respective status in "kubectl get pods".

```
→ cca2 git:(main) x kubectl get deployments
NAME                    READY   UP-TO-DATE   AVAILABLE   AGE
flask-app-deployment    2/2     2             2           118s
mongo-deployment        1/1     1             1           60s
→ cca2 git:(main) x kubectl get services
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP   PORT(S)          AGE
flask-app-service  NodePort    10.110.183.94 <none>        3000:30001/TCP   2m11s
kubernetes      ClusterIP   10.96.0.1     <none>        443/TCP          55m
mongo           ClusterIP   10.103.83.253 <none>        27017/TCP        73s
→ cca2 git:(main) x kubectl get pods
NAME                                                    READY   STATUS    RESTARTS   AGE
flask-app-deployment-b8cdcf6fb-dhpng                 1/1     Running   0           2m46s
flask-app-deployment-b8cdcf6fb-rjvhp                 1/1     Running   0           2m46s
mongo-deployment-cb654f8f-t84jl                      1/1     Running   0           108s
```

3) We run "minikube service flask-app-service" in order have minikube mount the IP and ports for local access. We are successfully given a URL with our live local deployment.

```

→ cca2 git:(main) ✗ minikube service flask-app-service

```

NAMESPACE	NAME	TARGET PORT	URL
default	flask-app-service	3000	http://192.168.49.2:30001

```

Starting tunnel for service flask-app-service.

```

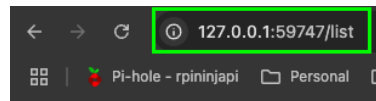
NAMESPACE	NAME	TARGET PORT	URL
default	flask-app-service		http://127.0.0.1:59747

```

Opening service default/flask-app-service in default browser...
Because you are using a Docker driver on darwin, the terminal needs to be open to run it.

```

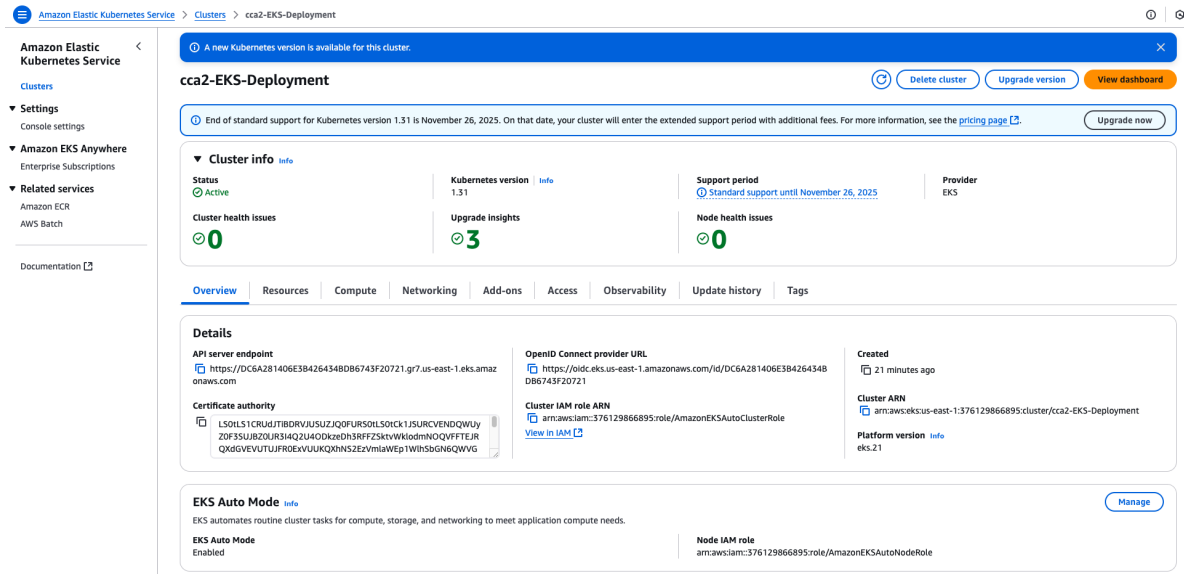
4) Finally, we navigate to the given URL and find that our application is deployed and working as expected.



ToDo Reminder

Part 4: Deploying the Application on AWS EKS

1) We create an AWS EKS Cluster via the AWS Management Console

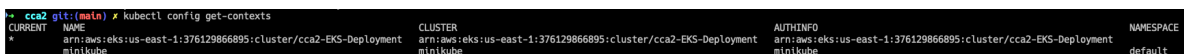


2) We then add the AWS EKS cluster to our kubeconfig

```
aws eks --region us-east-1 update-kubeconfig --cca2-EKS-Deployment
```

3) We switch context from minikube to AWS EKS

```
kubectl config use-context AWS_EKS
```



4) We apply deployment and services to AWS EKS Cluster

```
kubectl apply -f flask-eks-deployment.yaml
kubectl apply -f flask-eks-service.yaml
kubectl apply -f mongo-eks-deployment.yaml
kubectl apply -f mongo-eks-deployment.yaml
```

5) Use helm to setup AWS Load Balancer Controller

```
helm install aws-load-balancer-controller eks/aws-load-balancer-controller \
-n kube-system \
--set clusterName=cca2-EKS-Deployment \
--set serviceAccount.create=true \
--set serviceAccount.name=aws-load-balancer-controller \
--set region=us-east-1 \
--set vpcId=vpc-0fd462f3add9daa4b \
--set serviceAccount.annotations."eks\.amazonaws\.com/role arn"
    ="arn:aws:iam::376129866895:role/EKSLoadBalancerControllerRole" \
--set replicaCount=2
```

Note that the EKSLoadBalancerControllerRole, has a AWSLoadBalancerControllerIAMPolicy attached to it while also having a trust relationship configured with our AWS EKS cluster.

AWS Load Balancer Controller IAM Policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "iam:CreateServiceLinkedRole"
      ],
      "Resource": "*",
      "Condition": {
        "StringEquals": {
          "iam:AWSServiceName": "elasticloadbalancing.amazonaws.com"
        }
      }
    },
    {
      "Effect": "Allow",
      "Action": [
        "ec2:DescribeAccountAttributes",
        "ec2:DescribeAddresses",
        "ec2:DescribeAvailabilityZones",
        "ec2:DescribeInternetGateways",
        "ec2:DescribeVpcs",
        "ec2:DescribeVpcPeeringConnections",
        "ec2:DescribeSubnets",
        "ec2:DescribeSecurityGroups",
        "ec2:DescribeInstances",
        "ec2:DescribeNetworkInterfaces",
        "ec2:DescribeTags",
```

```

        "ec2:GetCoipPoolUsage",
        "ec2:DescribeCoipPools",
        "ec2:GetSecurityGroupsForVpc",
        "ec2:DescribeIpamPools",
        "elasticloadbalancing:DescribeLoadBalancers",
        "elasticloadbalancing:DescribeLoadBalancerAttributes",
        "elasticloadbalancing:DescribeListeners",
        "elasticloadbalancing:DescribeListenerCertificates",
        "elasticloadbalancing:DescribeSSLPolicies",
        "elasticloadbalancing:DescribeRules",
        "elasticloadbalancing:DescribeTargetGroups",
        "elasticloadbalancing:DescribeTargetGroupAttributes",
        "elasticloadbalancing:DescribeTargetHealth",
        "elasticloadbalancing:DescribeTags",
        "elasticloadbalancing:DescribeTrustStores",
        "elasticloadbalancing:DescribeListenerAttributes",
        "elasticloadbalancing:DescribeCapacityReservation"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "cognito-idp:DescribeUserPoolClient",
        "acm:ListCertificates",
        "acm:DescribeCertificate",
        "iam:ListServerCertificates",
        "iam:GetServerCertificate",
        "waf-regional:GetWebACL",
        "waf-regional:GetWebACLForResource",
        "waf-regional:AssociateWebACL",
        "waf-regional:DisassociateWebACL",
        "wafv2:GetWebACL",
        "wafv2:GetWebACLForResource",
        "wafv2:AssociateWebACL",
        "wafv2:DisassociateWebACL",
        "shield:GetSubscriptionState",
        "shield:DescribeProtection",
        "shield:CreateProtection",
        "shield>DeleteProtection"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:AuthorizeSecurityGroupIngress",

```



```

        "ec2:RevokeSecurityGroupIngress"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:CreateSecurityGroup"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:CreateTags"
    ],
    "Resource": "arn:aws:ec2:*:*:security-group/*",
    "Condition": {
        "StringEquals": {
            "ec2:CreateAction": "CreateSecurityGroup"
        },
        "Null": {
            "aws:RequestTag/elbv2.k8s.aws/cluster": "false"
        }
    }
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:CreateTags",
        "ec2>DeleteTags"
    ],
    "Resource": "arn:aws:ec2:*:*:security-group/*",
    "Condition": {
        "Null": {
            "aws:RequestTag/elbv2.k8s.aws/cluster": "true",
            "aws:ResourceTag/elbv2.k8s.aws/cluster": "false"
        }
    }
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:AuthorizeSecurityGroupIngress",
        "ec2:RevokeSecurityGroupIngress",
        "ec2>DeleteSecurityGroup"
    ],

```

```

    "Resource": "*",
    "Condition": {
        "Null": {
            "aws:ResourceTag/elbv2.k8s.aws/cluster": "false"
        }
    }
},
{
    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing:CreateLoadBalancer",
        "elasticloadbalancing:CreateTargetGroup"
    ],
    "Resource": "*",
    "Condition": {
        "Null": {
            "aws:RequestTag/elbv2.k8s.aws/cluster": "false"
        }
    }
},
{
    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing:CreateListener",
        "elasticloadbalancing>DeleteListener",
        "elasticloadbalancing:CreateRule",
        "elasticloadbalancing>DeleteRule"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing:AddTags",
        "elasticloadbalancing:RemoveTags"
    ],
    "Resource": [
        "arn:aws:elasticloadbalancing:*:*:targetgroup/*/*",
        "arn:aws:elasticloadbalancing:*:*:loadbalancer/net/*/*",
        "arn:aws:elasticloadbalancing:*:*:loadbalancer/app/*/*"
    ],
    "Condition": {
        "Null": {
            "aws:RequestTag/elbv2.k8s.aws/cluster": "true",
            "aws:ResourceTag/elbv2.k8s.aws/cluster": "false"
        }
    }
}

```

```

},
{
  "Effect": "Allow",
  "Action": [
    "elasticloadbalancing:AddTags",
    "elasticloadbalancing:RemoveTags"
  ],
  "Resource": [
    "arn:aws:elasticloadbalancing:*:listener/net/*/*/*",
    "arn:aws:elasticloadbalancing:*:listener/app/*/*/*",
    "arn:aws:elasticloadbalancing:*:listener-rule/net/*/*/*",
    "arn:aws:elasticloadbalancing:*:listener-rule/app/*/*/*"
  ]
},
{
  "Effect": "Allow",
  "Action": [
    "elasticloadbalancing:ModifyLoadBalancerAttributes",
    "elasticloadbalancing:SetIpAddressType",
    "elasticloadbalancing:SetSecurityGroups",
    "elasticloadbalancing:SetSubnets",
    "elasticloadbalancing>DeleteLoadBalancer",
    "elasticloadbalancing:ModifyTargetGroup",
    "elasticloadbalancing:ModifyTargetGroupAttributes",
    "elasticloadbalancing>DeleteTargetGroup",
    "elasticloadbalancing:ModifyListenerAttributes",
    "elasticloadbalancing:ModifyCapacityReservation",
    "elasticloadbalancing:ModifyIpPools"
  ],
  "Resource": "*",
  "Condition": {
    "Null": {
      "aws:ResourceTag/elbv2.k8s.aws/cluster": "false"
    }
  }
},
{
  "Effect": "Allow",
  "Action": [
    "elasticloadbalancing:AddTags"
  ],
  "Resource": [
    "arn:aws:elasticloadbalancing:*:targetgroup/*/*",
    "arn:aws:elasticloadbalancing:*:loadbalancer/net/*/*",
    "arn:aws:elasticloadbalancing:*:loadbalancer/app/*/*"
  ],
  "Condition": {

```

```

        "StringEquals": {
            "elasticloadbalancing:CreateAction": [
                "CreateTargetGroup",
                "CreateLoadBalancer"
            ]
        },
        "Null": {
            "aws:RequestTag/elbv2.k8s.aws/cluster": "false"
        }
    }
},
{
    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing:RegisterTargets",
        "elasticloadbalancing:DeregisterTargets"
    ],
    "Resource": "arn:aws:elasticloadbalancing:*:*:targetgroup/*/*"
},
{
    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing:SetWebAcl",
        "elasticloadbalancing:ModifyListener",
        "elasticloadbalancing:AddListenerCertificates",
        "elasticloadbalancing:RemoveListenerCertificates",
        "elasticloadbalancing:ModifyRule",
        "elasticloadbalancing:SetRulePriorities"
    ],
    "Resource": "*"
}
]
}

```

Trust Policy:

```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Principal": {
                "Federated": {
                    ↪ "arn:aws:iam::376129866895:oidc-provider/oidc.eks.us-east-1.amazonaws.com/id/DC6A281406E3B426434BDB6743F20721"
                }
            },
            "Action": "sts:AssumeRoleWithWebIdentity",

```

```

        "Condition": {
          "StringEquals": {
            "oidc.eks.us-east-1.amazonaws.com/id/DC6A281406E3B426434":
              ↪ BDB6743F20721:aud":
              ↪ "sts.amazonaws.com"
          }
        }
      }
    ]
  }
}

```

6) Tag one of the VPC Subnets being used by the AWS EKS Cluster so that it may be eligible for load balancer provisioning

Key: kubernetes.io/cluster/cca2-EKS-Deployment

Value: shared

Key: kubernetes.io/role/internal-elb

Value: 1

7) Check to see if the Load Balancer pods are running as expected

```

→ cca2 git:(main) ✗ kubectl get pods -n kube-system | grep aws-load-balancer-controller
aws-load-balancer-controller-8dd54586c-htjqp 1/1 Running 0 6h32m
aws-load-balancer-controller-8dd54586c-m9hbr 1/1 Running 0 6h32m

```

8) Check to see if deployment, services, and pods are running as expected

```

→ cca2 git:(main) ✗ kubectl get deployments
NAME                    READY    UP-TO-DATE    AVAILABLE    AGE
flask-app-deployment    2/2      2              2            8h
mongo-deployment        1/1      1              1            8h

```

```

→ cca2 git:(main) ✗ kubectl get pods
NAME                                READY    STATUS    RESTARTS    AGE
flask-app-deployment-8698687d8d-k6hj8 1/1      Running   10 (8h ago)  8h
flask-app-deployment-8698687d8d-wl8lg 1/1      Running   10 (8h ago)  8h
mongo-deployment-fb7c95b6f-b7psv      1/1      Running   0            8h

```

```

→ cca2 git:(main) ✗ kubectl get services
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP                                PORT(S)          AGE
flask-app-service LoadBalancer 10.100.84.34   k8s-default-flaskapp-2a98c1fa3b-7f5f467d2bc558b.elb.us-east-1.amazonaws.com 3000:32558/TCP   31m
kubernetes    ClusterIP   10.100.0.1    <none>                                     443/TCP          9h
mongo         ClusterIP   10.100.157.48 <none>                                     27017/TCP        8h

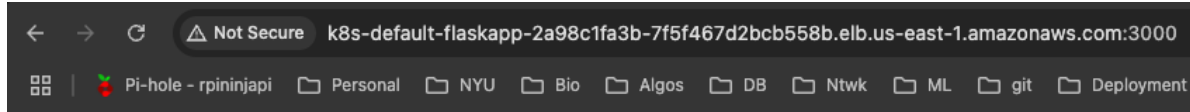
```

We can see that our services are all successfully deployed, our flask-app-service is assigned an external IP.

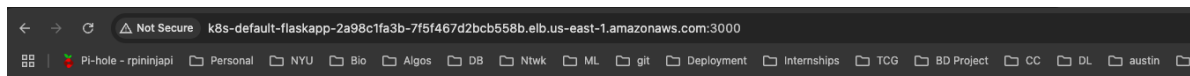
External IP:

http://k8s-default-flaskapp-2a98c1fa3b-7f5f467d2bcb558b.elb.us-east-1.amazonaws.com:3000/

9) We observe that our application is being deployed as expected on AWS EKS



ToDo Reminder



ToDo Reminder



Search Reference:

Unique ID

Search Task

To-Do LIST :

Status	Task Name	Description Name	Date
X	Aaron B.	gg	2025-03-22

Add a Task

Taskname

Enter Description here...

mm/dd/yyyy

Priority

Create

10) We show a final summary of the flask-app-service

```

caca2 git:(main) x kubectl describe service flask-app-service
Name: flask-app-service
Namespace: default
Labels: <none>
Annotations: service.beta.kubernetes.io/aws-load-balancer-internal: false
Selector: app=flask-app
Type: LoadBalancer
IP Family Policy: SingleStack
IP Families: IPv4
IP: 10.100.84.34
IPs: 10.100.84.34
LoadBalancer Ingress: k8s-default-flaskapp-2a98c1fa3b-7f5f467d2bcb558b.elb.us-east-1.amazonaws.com
Port: <unset> 3000/TCP
TargetPort: 3000/TCP
NodePort: <unset> 32558/TCP
Endpoints: 172.31.94.32:3000,172.31.94.33:3000
Session Affinity: None
External Traffic Policy: Cluster
Internal Traffic Policy: Cluster
Events:
  Type    Reason              Age   From      Message
  ----    -
  Normal  SuccessfullyReconciled 33m   service   Successfully reconciled

```

Part 5: Building the Replication Controller

1) We create a flask-rc.yaml file with the specifications of our replication controller, we will be starting with 2 replicas. We apply the replication controller yaml file using kubectl. Note, we are using "kind: Deployment" in our flask-eks-deployment.yaml file. The "Deployment" kind is a modern version of ReplicationController, with added features, in order to demonstrate that the replication controller is working as expected, we set the replica count in the flask-eks-deployment.yaml file to 0. Please reference the screenshots of the pods in sections 3 and 4 for additional context.

```
> cca2 git:(main) ✕ kubectl apply -f flask-eks-deployment.yaml
deployment.apps/flask-app-deployment configured
> cca2 git:(main) ✕ kubectl apply -f flask-rc.yaml
replicationcontroller/flask-app-rc unchanged
```

We print a summary of the pods after applying the changes in order to verify that the pod amount created matches our replica controllers replica count, we observe that this is true. The pod highlighted in green is to be deleted in the next step.

```
> cca2 git:(main) ✕ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
flask-app-rc-jz4qc	1/1	Running	0	23m
flask-app-rc-srwkl	1/1	Running	0	23m
mongo-deployment-fb7c95b6f-b7psv	1/1	Running	0	12h

2) We manually delete a pod in order to test our replica controller's response, the expected response is that a new flask-app pod will be created in place of the deleted flask-app pod. We observe that we obtain the expected response, a new flask-app pod is created.

```
> cca2 git:(main) ✕ kubectl delete pod flask-app-rc-jz4qc
pod "flask-app-rc-jz4qc" deleted
```

```
> cca2 git:(main) ✕ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
flask-app-rc-8kxjz	1/1	Running	0	5s
flask-app-rc-srwkl	1/1	Running	0	25m
mongo-deployment-fb7c95b6f-b7psv	1/1	Running	0	12h

3) We now scale up the replica count in our replica controller and apply the changes to our cluster

```
>→ cca2 git:(main) ✕ kubectl apply -f flask-rc.yaml
replicationcontroller/flask-app-rc configured
>→ cca2 git:(main) ✕ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
flask-app-rc-8kxjz	1/1	Running	0	32s
flask-app-rc-kptjq	1/1	Running	0	2s
flask-app-rc-srwkl	1/1	Running	0	25m
mongo-deployment-fb7c95b6f-b7psv	1/1	Running	0	12h

4) Finally, we scale our replica count in our replica controller back to the original 2 replicas, in order to demonstrate de-scaling. The application is still working and accessible as expected.

```
>→ cca2 git:(main) ✕ kubectl apply -f flask-rc.yaml
replicationcontroller/flask-app-rc configured
>→ cca2 git:(main) ✕ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
flask-app-rc-8kxjz	1/1	Running	0	54s
flask-app-rc-srwkl	1/1	Running	0	26m
mongo-deployment-fb7c95b6f-b7psv	1/1	Running	0	12h

Part 6: Applying a Rolling Update Strategy

1) We update our flask-eks-deployment.yaml file to include Strategy: RollingUpdate

```
! flask-eks-deployment.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: flask-app-deployment
5  spec:
6    replicas: 2
7    strategy:
8      type: RollingUpdate
9      rollingUpdate:
10       maxSurge: 1 # Allows up to x additional pods above desired replica count at time of update
11       maxUnavailable: 0 # Max unavailable pods during update
12  selector:
13    matchLabels:
14      app: flask-app
15  template:
16    metadata:
17      labels:
18        app: flask-app
19    spec:
20      containers:
21      - name: flask-app
22        image: aaronbengo/cca2:v2
23        ports:
24        - containerPort: 3000
25        env:
26        - name: FLASK_ENV
27          value: "development"
28        - name: MONGO_URI
29          value: "mongodb://mongo:27017/todo_db"
30        - name: PORT
31          value: "3000"
32
```

2) We build and push an updated docker image named aaronbengo/cca2:v2

```
➔ cca2 git:(main) ✗ docker buildx create --use
docker buildx build --platform linux/amd64,linux/arm64 \
  -t aaronbengo/cca2:v2 --push .

brave_kalam
[+] Building 30.7s (17/17) FINISHED
```

3) We check pods for status, apply the changes to our cluster, and recheck the status of the pods. We find that the pods are updated successfully.

```

→ cca2 git:(main) x kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-8698687d8d-m4cmm 1/1     Running   0           13m
flask-app-deployment-8698687d8d-xpv5k 1/1     Running   0           13m
mongo-deployment-fb7c95bbf-b/psv      1/1     Running   0           19h
→ cca2 git:(main) x kubectl apply -f flask-eks-deployment.yaml
deployment.apps/flask-app-deployment configured
→ cca2 git:(main) x kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-6544c6b4c5-n8j6w 1/1     Running   0           5s
flask-app-deployment-6544c6b4c5-sp4j5 1/1     Running   0           9s
mongo-deployment-fb7c95bbf-b/psv      1/1     Running   0           19h
→ cca2 git:(main) x kubectl rollout status deployment/flask-app-deployment
deployment "flask-app-deployment" successfully rolled out

```

4) We check the details of one of the new flask-app-deployment pods in order to verify the new image and check logs. We find that we indeed pulled and successfully updated the pods in the Events section.

```

→ cca2 git:(main) x kubectl describe pod flask-app-deployment-6544c6b4c5-n8j6w
Name: flask-app-deployment-6544c6b4c5-n8j6w
Namespace: default
Priority: 0
Service Account: default
Node: i-01dbade56c868046b/172.31.81.203
Start Time: Sun, 23 Mar 2025 08:04:44 -0400
Labels: app=flask-app
        pod-template-hash=6544c6b4c5
Annotations: <none>
Status: Running
IP: 172.31.94.35
IPs:
  IP: 172.31.94.35
Controlled By: ReplicaSet/flask-app-deployment-6544c6b4c5
Containers:
  flask-app:
    Container ID: containerd://dd75ce712357bbab0ea80f3778576d1f77425ccab2b3056c4c3058c7e6e46d9b
    Image: aaronbengo/cca2:v2
    Image ID: docker.io/aaronbengo/cca2@sha256:7eea4bd8fc1e08be214db50469adc84b88b0cb0f2aacf608e454e607b724088d
    Port: 3000/TCP
    Host Port: 0/TCP
    State: Running
      Started: Sun, 23 Mar 2025 08:04:45 -0400
    Ready: True
    Restart Count: 0
    Environment:
      FLASK_ENV: development
      MONGO_URI: mongod://mongo:27017/todo_db
      PORT: 3000
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-btwgj (ro)
Conditions:
  Type                               Status
  PodReadyToStartContainers          True
  Initialized                         True
  Ready                              True
  ContainersReady                    True
  PodScheduled                       True
Volumes:
  kube-api-access-btwgj:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
  QoS Class: BestEffort
  Node-Selectors: <none>
  Tolerations:
    node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
    node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type     Reason      Age    From          Message
  ----     -
  Normal   Scheduled   7m1s   default-scheduler   Successfully assigned default/flask-app-deployment-6544c6b4c5-n8j6w to i-01dbade56c868046b
  Normal   Pulling     7m     kubelet           Pulling image "aaronbengo/cca2:v2"
  Normal   Pulled      6m59s   kubelet           Successfully pulled image "aaronbengo/cca2:v2" in 701ms (701ms including waiting). Image size: 160395139 bytes.
  Normal   Created     6m59s   kubelet           Created container flask-app
  Normal   Started     6m59s   kubelet           Started container flask-app

```

5) We check the details of the flask deployment in order to verify that the updates occurred and that the new image version reflects the changes. We find that the updates were indeed successful

and that our deployment uses the new image named "aaronbengo/cca2:v2".

```
➤ cca2 git:(main) ✖ kubectl describe deployment flask-app-deployment
Name: flask-app-deployment
Namespace: default
CreationTimestamp: Sat, 22 Mar 2025 12:44:37 -0400
Labels: <none>
Annotations: deployment.kubernetes.io/revision: 2
Selector: app=flask-app
Replicas: 2 desired | 2 updated | 2 total | 2 available | 0 unavailable
StrategyType: RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 0 max unavailable, 1 max surge
Pod Template:
  Labels: app=flask-app
  Containers:
    flask-app:
      Image: aaronbengo/cca2:v2
      Port: 3000/TCP
      Host Port: 0/TCP
      Environment:
        FLASK_ENV: development
        MONGO_URI: mongodb://mongo:27017/todo_db
        PORT: 3000
      Mounts: <none>
      Volumes: <none>
      Node-Selectors: <none>
      Tolerations: <none>
  Conditions:
    Type      Status  Reason
    ----      -
    Available  True    MinimumReplicasAvailable
    Progressing True    NewReplicaSetAvailable
  OldReplicaSets: flask-app-deployment-8698687d8d (0/0 replicas created)
  NewReplicaSet: flask-app-deployment-6544c6b4c5 (2/2 replicas created)
  Events:
    Type      Reason      Age    From      Message
    ----      -
    Normal    ScalingReplicaSet  38m    deployment-controller    Scaled up replica set flask-app-deployment-8698687d8d to 2 from 0
    Normal    ScalingReplicaSet  24m    deployment-controller    Scaled up replica set flask-app-deployment-6544c6b4c5 to 1
    Normal    ScalingReplicaSet  24m    deployment-controller    Scaled down replica set flask-app-deployment-8698687d8d to 1 from 2
    Normal    ScalingReplicaSet  24m    deployment-controller    Scaled up replica set flask-app-deployment-6544c6b4c5 to 2 from 1
    Normal    ScalingReplicaSet  24m    deployment-controller    Scaled down replica set flask-app-deployment-8698687d8d to 0 from 1
```

Part 7: Applying a Health Monitoring Strategy

- 1) We add an API endpoint in order to allow for health checks

```
@app.route("/health")
def hello ():
    return "OK", 200
```

- 2) We build and push the new image named "aaronbengo/cca2:v3"

```
➤ cca2 git:(main) ✗ docker buildx create --use
docker buildx build --platform linux/amd64,linux/arm64 \
  -t aaronbengo/cca2:v3 --push .

elated_hamilton
[+] Building 28.8s (17/17) FINISHED
```

- 3) We update our flask-eks-development.yaml file to include the livenessProbe and readinessProbe. Our probes check our app's health by performing an HTTP get to our /health endpoint.

```

! flask-eks-deployment.yaml
1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: flask-app-deployment
5  spec:
6    replicas: 2
7    strategy:
8      type: RollingUpdate
9      rollingUpdate:
10       maxSurge: 1 # Allows up to x additional pods above desired replica count at time of update
11       maxUnavailable: 0 # Max unavailable pods during update
12    selector:
13      matchLabels:
14        app: flask-app
15    template:
16      metadata:
17        labels:
18          app: flask-app
19      spec:
20        containers:
21          - name: flask-app
22            image: aaronbengo/cca2:v3
23            ports:
24              - containerPort: 3000
25            env:
26              - name: FLASK_ENV
27                value: "development"
28              - name: MONGO_URI
29                value: "mongodb://mongo:27017/todo_db"
30              - name: PORT
31                value: "3000"
32            livenessProbe:
33              httpGet:
34                path: /health # Uses HTTP get to check health status of app
35                port: 3000
36              initialDelaySeconds: 10 # Performs the first check after 10s have passed from container creation
37              periodSeconds: 10 # Checks every 10 seconds
38              failureThreshold: 3 # Amount of consecutive failures before k8 restarts the pod
39            readinessProbe:
40              httpGet:
41                path: /health # Uses HTTP get to check health status of app
42                port: 3000
43              initialDelaySeconds: 5 # Performs the first check after 5s have passed from container creation
44              periodSeconds: 5 # Checks every 5 seconds
45              failureThreshold: 2 # Amount of consecutive failures before k8 restarts the pod
46

```

4) We apply the changes to our cluster and we observe that new pods are spun up and used to replace the old pods.

```

→ cca2 git:(main) ✖ kubectl apply -f flask-eks-deployment.yaml
deployment.apps/flask-app-deployment configured
→ cca2 git:(main) ✖ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-6544c6b4c5-n8j6w 1/1     Running   0           88m
flask-app-deployment-6544c6b4c5-sp4j5 1/1     Running   0           88m
flask-app-deployment-6cdfb5b4c5-fnr2x 0/1     Running   0           6s
mongo-deployment-fb7c95b6f-b7psv      1/1     Running   0           20h
→ cca2 git:(main) ✖ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-6544c6b4c5-sp4j5 1/1     Running   0           89m
flask-app-deployment-6cdfb5b4c5-dqdrd 0/1     Running   0           8s
flask-app-deployment-6cdfb5b4c5-fnr2x 1/1     Running   0           19s
mongo-deployment-fb7c95b6f-b7psv      1/1     Running   0           20h
→ cca2 git:(main) ✖ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-6cdfb5b4c5-dqdrd 1/1     Running   0           21s
flask-app-deployment-6cdfb5b4c5-fnr2x 1/1     Running   0           32s
mongo-deployment-fb7c95b6f-b7psv      1/1     Running   0           20h
→ cca2 git:(main) ✖ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-6cdfb5b4c5-dqdrd 1/1     Running   0           38s
flask-app-deployment-6cdfb5b4c5-fnr2x 1/1     Running   0           49s
mongo-deployment-fb7c95b6f-b7psv      1/1     Running   0           20h

```

5) We check the description of the new pods for Liveness and Readiness metadata fields, we confirm that the new pods are using the appropriate image.

```

+ cca2 git:(main) x kubectl describe pod flask-app-deployment-6cdfb5b4c5-dqdrd
Name: flask-app-deployment-6cdfb5b4c5-dqdrd
Namespace: default
Priority: 0
Service Account: default
Node: i-01dbade56c868046b/172.31.81.203
Start Time: Sun, 23 Mar 2025 09:33:40 -0400
Labels: app=flask-app
        pod-template-hash=6cdfb5b4c5
Annotations: <none>
Status: Running
IP: 172.31.94.33
IPs:
  IP: 172.31.94.33
Controlled By: ReplicaSet/flask-app-deployment-6cdfb5b4c5
Containers:
  flask-app:
    Container ID: containerd://ebc2a482dba079ebe56ff8b92e5617df297c9a2bc9c7f8e44be79e8020e8b4f1
    Image: aaronbengo/cca2:v3
    Image ID: docker.io/aaronbengo/cca2@sha256:d773af01ec73aaae4a7de51a65c1c4511b4a621284fd06df983a09fe70e16a03
    Port: 3000/TCP
    Host Port: 0/TCP
    State: Running
      Started: Sun, 23 Mar 2025 09:33:41 -0400
    Ready: True
    Restart Count: 0
    Liveness: http-get http://:3000/health delay=10s timeout=1s period=10s #success=1 #failure=3
    Readiness: http-get http://:3000/health delay=5s timeout=1s period=5s #success=1 #failure=2
    Environment:
      FLASK_ENV: development
      MONGO_URI: mongodb://mongo:27017/todo_db
      PORT: 3000
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-l7z6l (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized         True
  Ready               True
  ContainersReady     True
  PodScheduled        True
Volumes:
  kube-api-access-l7z6l:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type    Reason      Age   From          Message
  ----    -
  Normal  Scheduled   72s   default-scheduler  Successfully assigned default/flask-app-deployment-6cdfb5b4c5-dqdrd to i-01dbade56c868046b
  Normal  Pulling     72s   kubelet        Pulling image "aaronbengo/cca2:v3"
  Normal  Pulled      71s   kubelet        Successfully pulled image "aaronbengo/cca2:v3" in 118ms (118ms including waiting). Image size: 160395290 bytes.
  Normal  Created     71s   kubelet        Created container flask-app
  Normal  Started     71s   kubelet        Started container flask-app

```

6) Now we simulate an error by returning an error flag in the /health api response

```

@app.route("/health")
def hello ():
    return "Error", 500

```

7) We rebuild the image with the /health api update, we name it "aaronbengo/cca2:v4" and we push it to DockerHub

```

➔ cca2 git:(main) x docker buildx create --use
docker buildx build --platform linux/amd64,linux/arm64 \
  -t aaronbengo/cca2:v4 --push .

elastic_clarke
[+] Building 34.6s (17/17) FINISHED

```


8) We apply the changes to our EKS cluster and monitor our pods, we see that a new pod is created but not yet ready. We check the pods description for image name, Liveness, Readiness, and Event logs. We find that the service status is switched to Unhealthy due to the 500 statuscode responses, both the Liveness and Readiness probes fail, which prompts the system to try and restart the pod. Since our system continually gets 500 status codes, the pod never actually mounts, we can assume that if the 500 status codes became 200 at one point, then the pod would successfully mount instead of attempting to restart.

```

+ cca2 git:(main) x kubectl describe pod flask-app-deployment-698dd4766-5dzdc
Name: flask-app-deployment-698dd4766-5dzdc
Namespace: default
Priority: 0
Service Account: default
Node: i-01dbade56c868046b/172.31.81.203
Start Time: Sun, 23 Mar 2025 10:01:05 -0400
Labels: app=flask-app
        pod-template-hash=698dd4766
Annotations: <none>
Status: Running
IP: 172.31.94.35
IPs:
  IP: 172.31.94.35
Controlled By: ReplicaSet/flask-app-deployment-698dd4766
Containers:
  flask-app:
    Container ID: containerd://ddacc2938326df5e012304b01c20d6007fda8a9f4f8bbd935716e97fd0ba5c9f
    Image: aaronbengo/cca2:v4
    Image ID: docker.io/aaronbengo/cca2@sha256:4a00e80c181731ed56eb0f68cafec0d82684512a37de741d8a50d637b93940b
    Port: 3000/TCP
    Host Port: 0/TCP
    State: Waiting
      Reason: CrashLoopBackOff
    Last State: Terminated
      Reason: Completed
      Exit Code: 0
      Started: Sun, 23 Mar 2025 10:03:45 -0400
      Finished: Sun, 23 Mar 2025 10:04:15 -0400
    Ready: False
    Restart Count: 4
    Liveness: http-get http://:3000/health delay=10s timeout=1s period=10s #success=1 #failure=3
    Readiness: http-get http://:3000/health delay=5s timeout=1s period=5s #success=1 #failure=2
    Environment:
      FLASK_ENV: development
      MONGO_URI: mongodb://mongo:27017/todo_db
      PORT: 3000
    Mounts:
      /var/run/secrets/kubernetes.io/serviceaccount from kube-api-access-59vch (ro)
Conditions:
  Type              Status
  PodReadyToStartContainers  True
  Initialized         True
  Ready               False
  ContainersReady      False
  PodScheduled        True
Volumes:
  kube-api-access-59vch:
    Type: Projected (a volume that contains injected data from multiple sources)
    TokenExpirationSeconds: 3607
    ConfigMapName: kube-root-ca.crt
    ConfigMapOptional: <nil>
    DownwardAPI: true
QoS Class: BestEffort
Node-Selectors: <none>
Tolerations: node.kubernetes.io/not-ready:NoExecute op=Exists for 300s
              node.kubernetes.io/unreachable:NoExecute op=Exists for 300s
Events:
  Type      Reason      Age          From          Message
  ----      -
  Normal    Scheduled   3m23s        default-scheduler   Successfully assigned default/flask-app-deployment-698dd4766-5dzdc to i-01dbade56c868046b
  Normal    Pulled      3m20s        kubelet         Successfully pulled image "aaronbengo/cca2:v4" in 3.219s (3.219s including waiting). Image size: 160395834 bytes.
  Normal    Pulling     2m43s (x2 over 3m23s)  kubelet         Pulling image "aaronbengo/cca2:v4"
  Normal    Created     2m43s (x2 over 3m20s)  kubelet         Created container flask-app
  Normal    Killing     2m43s        kubelet         Container flask-app failed liveness probe, will be restarted
  Normal    Pulled      2m43s        kubelet         Successfully pulled image "aaronbengo/cca2:v4" in 114ms (114ms including waiting). Image size: 160395834 bytes.
  Normal    Started     2m42s (x2 over 3m20s)  kubelet         Started container flask-app
  Warning   Unhealthy   2m23s (x4 over 3m3s)  kubelet         Liveness probe failed: HTTP probe failed with statuscode: 500
  Warning   Unhealthy   2m18s (x12 over 3m13s) kubelet         Readiness probe failed: HTTP probe failed with statuscode: 500

```

9) We conclude by resetting the /health api to return a status code of 200 so that the livenessProbe and readinessProbe can work as expected. With this approach, our system will be able to discern the health of our application tactically. We also confirm that our deployment is still working as expected, and our webpage is live.

```

+ cca2 git:(main) x kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
flask-app-deployment-6445687dcc-9dq7r 1/1     Running   0           110s
flask-app-deployment-6445687dcc-9n4d8 1/1     Running   0           2m50s
mongo-deployment-fb7c95b6f-b7p5v      1/1     Running   0           21h

```

We show the logs of one the new pods showing it is checking the /health api endpoint for a 200 status code.

```
→ cca2 git:(main) ✕ kubectl logs flask-app-deployment-6445687dcc-9n4d8
* Serving Flask app 'app.py' (lazy loading)
* Environment: development
* Debug mode: on
* Running on all addresses.
  WARNING: This is a development server. Do not use it in a production deployment.
* Running on http://172.31.94.36:3000/ (Press CTRL+C to quit)
* Restarting with stat
* Debugger is active!
* Debugger PIN: 527-281-284
172.31.81.203 - - [23/Mar/2025 14:23:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:23:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:24:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:24:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:25:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:25:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:26:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:26:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:27:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:27:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:28:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:28:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:29:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:29:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:30:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:30:18] "GET /health HTTP/1.1" 200 -
172.31.22.84 - - [23/Mar/2025 14:30:34] "GET /static/assets/style.css HTTP/1.1" 304 -
172.31.22.84 - - [23/Mar/2025 14:30:34] "GET /static/images/yes.png HTTP/1.1" 304 -
172.31.22.84 - - [23/Mar/2025 14:30:34] "GET /static/images/no.png HTTP/1.1" 304 -
172.31.81.203 - - [23/Mar/2025 14:31:18] "GET /health HTTP/1.1" 200 -
172.31.81.203 - - [23/Mar/2025 14:31:18] "GET /health HTTP/1.1" 200 -
```

Part 7: Alerting with Prometheus

- 1) First we install Prometheus onto our cluster

```
➔ cca2 git:(main) x helm install prometheus prometheus-community/kube-prometheus-stack \
--namespace monitoring --create-namespace

NAME: prometheus
LAST DEPLOYED: Sun Mar 23 17:52:31 2025
NAMESPACE: monitoring
STATUS: deployed
REVISION: 1
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace monitoring get pods -l "release=prometheus"

Get Grafana 'admin' user password by running:

  kubectl --namespace monitoring get secrets prometheus-grafana -o jsonpath="{.data.admin-password}" | base64 -d ; echo

Access Grafana local instance:

  export POD_NAME=$(kubectl --namespace monitoring get pod -l "app.kubernetes.io/name=grafana,app.kubernetes.io/instance=prometheus" -oname)
  kubectl --namespace monitoring port-forward $POD_NAME 3000

Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager and Prometheus instances using the Operator.
```

- 2) We then check if the Prometheus and Alertmanager pods were set up successfully

```
➔ cca2 git:(main) x kubectl get pods -n monitoring
```

NAME	READY	STATUS	RESTARTS	AGE
alertmanager-prometheus-kube-prometheus-alertmanager-0	2/2	Running	0	41s
prometheus-grafana-75bb7d6986-pvlkc	3/3	Running	0	46s
prometheus-kube-prometheus-operator-65c669f8f9-95hfk	1/1	Running	0	46s
prometheus-kube-state-metrics-645c667b6-xqrk	1/1	Running	0	46s
prometheus-prometheus-kube-prometheus-prometheus-0	2/2	Running	0	41s
prometheus-prometheus-node-exporter-42cp9	1/1	Running	0	46s
prometheus-prometheus-node-exporter-4rjtb	1/1	Running	0	46s

- 3) We check the active rules in the prometheusrule monitoring namespace

```

→ cca2 git:(main) ✗ kubectl get prometheusrule -n monitoring

```

NAME	AGE
prometheus-kube-prometheus-alertmanager.rules	16s
prometheus-kube-prometheus-config-reloaders	16s
prometheus-kube-prometheus-etcd	16s
prometheus-kube-prometheus-general.rules	16s
prometheus-kube-prometheus-k8s.rules.container-cpu-usage-second	16s
prometheus-kube-prometheus-k8s.rules.container-memory-cache	16s
prometheus-kube-prometheus-k8s.rules.container-memory-rss	16s
prometheus-kube-prometheus-k8s.rules.container-memory-swap	16s
prometheus-kube-prometheus-k8s.rules.container-memory-working-s	16s
prometheus-kube-prometheus-k8s.rules.container-resource	16s
prometheus-kube-prometheus-k8s.rules.pod-owner	16s
prometheus-kube-prometheus-kube-apiserver-availability.rules	16s
prometheus-kube-prometheus-kube-apiserver-burnrate.rules	16s
prometheus-kube-prometheus-kube-apiserver-histogram.rules	16s
prometheus-kube-prometheus-kube-apiserver-slos	16s
prometheus-kube-prometheus-kube-prometheus-general.rules	16s
prometheus-kube-prometheus-kube-prometheus-node-recording.rules	16s
prometheus-kube-prometheus-kube-scheduler.rules	16s
prometheus-kube-prometheus-kube-state-metrics	16s
prometheus-kube-prometheus-kubelet.rules	16s
prometheus-kube-prometheus-kubernetes-apps	16s
prometheus-kube-prometheus-kubernetes-resources	16s
prometheus-kube-prometheus-kubernetes-storage	16s
prometheus-kube-prometheus-kubernetes-system	16s
prometheus-kube-prometheus-kubernetes-system-apiserver	16s
prometheus-kube-prometheus-kubernetes-system-controller-manager	16s
prometheus-kube-prometheus-kubernetes-system-kube-proxy	16s
prometheus-kube-prometheus-kubernetes-system-kubelet	16s
prometheus-kube-prometheus-kubernetes-system-scheduler	16s
prometheus-kube-prometheus-node-exporter	16s
prometheus-kube-prometheus-node-exporter.rules	16s
prometheus-kube-prometheus-node-network	16s
prometheus-kube-prometheus-node.rules	16s
prometheus-kube-prometheus-prometheus	16s
prometheus-kube-prometheus-prometheus-operator	16s

4) We setup a custom-alerts.yaml manifest for our PrometheusRule rules

```

1  ! custom-alerts.yaml
2  apiVersion: monitoring.coreos.com/v1
3  kind: PrometheusRule
4  metadata:
5    name: custom-alerts
6    namespace: monitoring
7    labels:
8      release: prometheus # This label matches the ruleSelector of your Prometheus instance.
9      role: alert-rules # Optionally, you can keep additional labels if desired.
10 spec:
11   groups:
12     - name: alerts
13       rules:
14         - alert: PodDeletionDetectedForFlaskDeployment
15           expr: kube_deployment_status_replicas_available{deployment="flask-app-deployment", namespace="default"}
16                 < kube_deployment_spec_replicas{deployment="flask-app-deployment", namespace="default"}
17           for: 1m
18           labels:
19             severity: critical
20           annotations:
21             summary: "Pod deletion detected for deployment {{ $labels.deployment }}"
22             description: "The deployment {{ $labels.deployment }} is missing pods: desired replicas are less than available replica
23         - alert: PodCrashLooping
24           expr: kube_pod_container_status_restarts_total{namespace="default", container="flask-app"} > 1
25           for: 2m
26           labels:
27             severity: critical
28           annotations:
29             summary: "Pod {{ $labels.pod }} in namespace {{ $labels.namespace }} is crash looping"
30             description: "The pod has restarted more than 1 time in the last 2 minutes."
31         - alert: DeploymentHasTooFewReplicas
32           expr: kube_deployment_status_replicas_unavailable{deployment="flask-app-deployment", namespace="default"} > 0
33           for: 2m
34           labels:
35             severity: warning
36           annotations:
37             summary: "Deployment {{ $labels.deployment }} has unavailable replicas"
38             description: "The deployment does not have all replicas running."

```

5) We apply the custom-alerts manifest to our cluster

```

➔ cca2 git:(main) x kubectl apply -f custom-alerts.yaml -n monitoring
prometheusrule.monitoring.coreos.com/custom-alerts created

```

6) Again we check the active rules in the prometheusrule monitoring namespace

```

➔ cca2 git:(main) ✕ kubectl get prometheusrule -n monitoring

```

NAME	AGE
custom-alerts	35s
prometheus-kube-prometheus-alertmanager.rules	4m5s
prometheus-kube-prometheus-config-reloaders	4m5s
prometheus-kube-prometheus-etcd	4m5s
prometheus-kube-prometheus-general.rules	4m5s
prometheus-kube-prometheus-k8s.rules.container-cpu-usage-second	4m5s
prometheus-kube-prometheus-k8s.rules.container-memory-cache	4m5s
prometheus-kube-prometheus-k8s.rules.container-memory-rss	4m5s
prometheus-kube-prometheus-k8s.rules.container-memory-swap	4m5s
prometheus-kube-prometheus-k8s.rules.container-memory-working-s	4m5s
prometheus-kube-prometheus-k8s.rules.container-resource	4m5s
prometheus-kube-prometheus-k8s.rules.pod-owner	4m5s
prometheus-kube-prometheus-kube-apiserver-availability.rules	4m5s
prometheus-kube-prometheus-kube-apiserver-burnrate.rules	4m5s
prometheus-kube-prometheus-kube-apiserver-histogram.rules	4m5s
prometheus-kube-prometheus-kube-apiserver-slos	4m5s
prometheus-kube-prometheus-kube-prometheus-general.rules	4m5s
prometheus-kube-prometheus-kube-prometheus-node-recording.rules	4m5s
prometheus-kube-prometheus-kube-scheduler.rules	4m5s
prometheus-kube-prometheus-kube-state-metrics	4m5s
prometheus-kube-prometheus-kubelet.rules	4m5s
prometheus-kube-prometheus-kubernetes-apps	4m5s
prometheus-kube-prometheus-kubernetes-resources	4m5s
prometheus-kube-prometheus-kubernetes-storage	4m5s
prometheus-kube-prometheus-kubernetes-system	4m5s
prometheus-kube-prometheus-kubernetes-system-apiserver	4m5s
prometheus-kube-prometheus-kubernetes-system-controller-manager	4m5s
prometheus-kube-prometheus-kubernetes-system-kube-proxy	4m5s
prometheus-kube-prometheus-kubernetes-system-kubelet	4m5s
prometheus-kube-prometheus-kubernetes-system-scheduler	4m5s
prometheus-kube-prometheus-node-exporter	4m5s
prometheus-kube-prometheus-node-exporter.rules	4m5s
prometheus-kube-prometheus-node-network	4m5s
prometheus-kube-prometheus-node.rules	4m5s
prometheus-kube-prometheus-prometheus	4m5s
prometheus-kube-prometheus-prometheus-operator	4m5s

7) We show a snapshot of the prometheus UI's alert tab, which shows our three custom rules active, this means that prometheus was able to sync them using the prometheus operator.



Filter by rule state

Filter by rule name or labels

<

1

2

>

alerts /etc/prometheus/rules/prometheus-prometheus-kube-prometheus-prometheus-rulefi

PodDeletionDetectedForFlaskDeployment

PodCrashLooping

DeploymentHasTooFewReplicas

alertmanager.rules /etc/prometheus/rules/prometheus-prometheus-kube-prometheus-

AlertmanagerFailedReload

AlertmanagerMembersInconsistent

AlertmanagerFailedToSendAlerts

AlertmanagerClusterFailedToSendAlerts

AlertmanagerClusterFailedToSendAlerts

AlertmanagerConfigInconsistent

AlertmanagerClusterDown

AlertmanagerClusterCrashlooping

8) Next, we shift our focus to slack messages on the event that one of our alerts is triggered. We create the alertmanager-override.yaml manifest in order to specify the alertmanager config.

```
! alertmanager-overrides.yaml
1 alertmanager:
2   config:
3     global:
4       slack_api_url: "https://hooks.slack.com/services/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"
5     route:
6       group_by: ['alertname']           # grouping keys (adjust as needed)
7       group_wait: 30s
8       group_interval: 5m
9       repeat_interval: 12h
10      receiver: slack-notifications     # send all alerts to Slack by default
11      routes:
12        - matchers:
13          - alertname = "Watchdog"      # drop the Watchdog heartbeat alert
14          receiver: null
15      receivers:
16        - name: "null"
17        - name: "slack-notifications"
18          slack_configs:
19            - channel: "#test"           # Slack channel for alerts
20              send_resolved: true
21              # Optional: title, text, etc., or rely on Alertmanager defaults/templates
22
```

9) We use helm in order to upgrade the prometheus deployment, this time we pass our alertmanager-override.yaml manifest as an argument for the alertmanager service config.

```
→ cca2 git:(main) x helm upgrade prometheus prometheus-community/kube-prometheus-stack \
--namespace monitoring \
-f alertmanager-overrides.yaml

Release "prometheus" has been upgraded. Happy Helming!
NAME: prometheus
LAST DEPLOYED: Sun Mar 23 19:14:47 2025
NAMESPACE: monitoring
STATUS: deployed
REVISION: 3
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace monitoring get pods -l "release=prometheus"

Get Grafana 'admin' user password by running:

  kubectl --namespace monitoring get secrets prometheus-grafana -o jsonpath="{.data.admin-password}" | base64 -d ; echo

Access Grafana local instance:

  export POD_NAME=$(kubectl --namespace monitoring get pod -l "app.kubernetes.io/name=grafana,app.kubernetes.io/instance=prometheus" -oname)
  kubectl --namespace monitoring port-forward $POD_NAME 3000

Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager and Prometheus instances using the Operator.
```

10) We observe that the receiver is now live on the Alertmanager UI. It is important to note that my submission lacks my 3 defined custom alerts. I had issues trying to get them added and was unable to in by the deadline of the assignment, I am sharing all the progress that I was able to make.

Alertmanager

Alerts

Silences

Status

Settings

Help

Filter

Group

Custom matcher, e.g. `env="production"`

+ Expand all groups

+ slack-notifications

+ slack-notifications

+ slack-notifications

+ slack-notifications

alertname="KubeControllerManagerDown" + 1 alert

alertname="KubeProxyDown" + 1 alert

alertname="KubeSchedulerDown" + 1 alert

alertname="Watchdog" + 1 alert

11) I conclude by showing that the services that appear in the Alertmanager UI were live and ready to communicate with my slack channel.

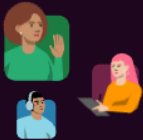
test

Messages +

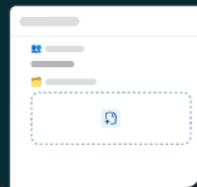
👋 Welcome to the # test channel

This channel is for everything # test. Get started by setting up the channel for your team:

Invite teammates
Add your whole team



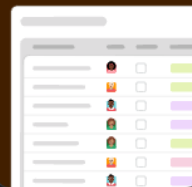
Add project brief
Canvas template



Host weekly syncs
Huddle in Slack



Add project tracker
List template



Today ▾



ab6503 11:04 AM

added an integration to this channel: [incoming-webhook](#)



Alertmanager APP 7:28 PM

[FIRING:1] Watchdog (monitoring/prometheus-kube-prometheus-prometheus none)

[FIRING:1] KubeSchedulerDown (monitoring/prometheus-kube-prometheus-prometheus critical)

[FIRING:1] KubeControllerManagerDown (monitoring/prometheus-kube-prometheus-prometheus critical)

[FIRING:1] KubeProxyDown (monitoring/prometheus-kube-prometheus-prometheus critical)

B I

Message #test

+ Aa